

UMAR FAROOQUE

☎ +91 6387273062 | ✉ umar80643@gmail.com | 🔗 LinkedIn -<https://www.linkedin.com/in/umar---farooque/> | 🐙 GitHub-<https://github.com/umar80643/> | 🌐 Portfolio- <https://umar-portfolio-bay.vercel.app/> | 📊 LeetCode:200+Problems Solved

PROFESSIONAL SUMMARY

AI Engineer and Computer Science undergraduate experienced in building multi-agent AI systems, LLM-integrated applications, and evidence-grounded RAG workflows using Python, LangGraph, and Agno, with production-style FastAPI/Docker services and hands-on model orchestration across Gemini, OpenAI, and Ollama.

TECHNICAL SKILLS

Languages: Python, Java, SQL, JavaScript, HTML/CSS

Generative AI & LLMs: Google Gemini, OpenAI, Ollama, Prompt Engineering, LLM Integration

AI Agents & Orchestration: LangGraph, LangChain, Agno, Multi-Agent Systems, MCP-Style Tool Boundaries, Agentic Workflows

Retrieval & Context Engineering: Hybrid RAG (Lexical + Vector, RRF), FAISS, Qdrant, Evidence Grounding, Source-Quality Classification

Backend & APIs: FastAPI, REST APIs, SQLAlchemy, AioHTTP, API-Key Auth, Rate Limiting

Databases: PostgreSQL, Microsoft SQL Server, SQLite, Redis

DevOps & Testing: Docker, Docker Compose, GitHub Actions CI/CD, Git, Pytest, Prometheus, Grafana, k6

Core CS: DSA, OOP & SOLID, DBMS, Async Programming, System Architecture, NLP

EXPERIENCE

Data Analytics Job Simulation — Deloitte (via Forage)

August 2026

- Completed a Forage-hosted virtual simulation applying structured data-analysis methods to a forensic technology case scenario, identifying patterns in a dataset and translating findings into a decision-ready summary.
- Practiced communicating technical analysis to a non-technical audience per Deloitte's professional consulting standards.

PROJECTS

Enterprise AI Software Engineering & Incident Resolution Agent | Python, FastAPI, LangGraph, LangChain, Qdrant, PostgreSQL, Redis, Next.js, Docker

- Architected an evidence-grounded incident-investigation system orchestrated with LangGraph, coordinating planner, repository-analysis, retrieval, debugging, root-cause, solution, and reviewer agents through a typed, resumable state machine with bounded retries.
- Implemented hybrid lexical + vector retrieval with reciprocal-rank fusion over an indexed codebase, validating citations against retrieved chunks and returning INSUFFICIENT_EVIDENCE instead of unsupported conclusions.
- Designed MCP-style permissioned tool boundaries (read-only, safe-execution, write-repository, create-PR) with mandatory human approval before repository writes, treating repository content as untrusted evidence to prevent prompt injection.

GitHub- <https://github.com/umar80643/enterprise-ai-incident-response-agent.git>

Product Launch Intelligence Agent | Python, Agno, OpenAI, Firecrawl, FastAPI, Streamlit, SQLAlchemy, Docker

- Architected a multi-agent workflow with 3 Agno specialist agents for competitor analysis, market sentiment, and launch metrics, using Firecrawl-backed web research to generate evidence-tagged intelligence reports.
- Built a layered FastAPI service with API-key auth, rate limiting, retries, and SQLAlchemy persistence; validated with 22 pytest tests and GitHub Actions CI/CD.

GitHub - <https://github.com/umar80643/PRODUCT-LAUNCH-INTELLIGENCE-AGENT.git>

Plagiarism & AI-Content Detector | Python, FastAPI, Streamlit, FAISS, Scikit-learn, Transformers, Docker

- Built a production-style plagiarism and AI-text detection application with FastAPI, API-key authentication, rate limiting, a decoupled Streamlit client, and Docker Compose service orchestration.
- Designed configurable retrieval with TF-IDF+LSA or Sentence-Transformers embeddings, FAISS vector search, optional cross-encoder reranking, and a lexical fallback.
- Developed a 3-tier AI-text detector spanning heuristic, trained Scikit-learn, and fine-tuned Transformer modes, with incremental corpus indexing and calibrated confidence scoring.
- Trained a TF-IDF + Logistic Regression classifier on a held-out split and evaluated accuracy, precision, recall, F1, and ROC-AUC; independently unit-tested temperature-scaling calibration. **GitHub-** <https://github.com/umar80643/plagiarism-ai-detector.git>

L7 Load Balancer | Python, AsyncIO, AioHTTP, Docker, Prometheus, Grafana, k6

- Built a Layer 7 HTTP load balancer with AsyncIO/AioHTTP, sustaining ~2,009 req/s at 0% error rate and 35.42 ms p95 latency in k6 peak-throughput benchmarks.
- Designed 6 pluggable routing algorithms with active/passive health checks, circuit breaking, exponential-backoff retries, token-bucket rate limiting, and sticky sessions.
- Wrote 133 unit and real-network integration tests with 97% line coverage and a GitHub Actions CI pipeline covering linting, tests, Docker build, and Docker Compose smoke testing. **GitHub -** <https://github.com/umar80643/L7-load-balancer.git>

EDUCATION

AllenHouse Institute of Technology, Kanpur, UP — Bachelor of Technology, Computer Science

CGPA: 7.2 | 2023 – 2027 (Expected)

UP Kirana Seva Samiti Vidyalaya, Kanpur, UP — Intermediate (Class XII), 2022

TRAINING & CERTIFICATIONS

- Java Certification — Scaler.com (2025): Object-Oriented Programming, multithreading, Collections Framework, and core data structures.
- Python Certification — Udemy: core language fundamentals, file I/O, REST API integration, and data manipulation with pandas & NumPy.
- Data Structures & Algorithms Training — AllenHouse Institute of Technology (2025): algorithm design and complexity analysis across Arrays, Trees, Graphs, and Dynamic Programming.